

Nam Dao

Toronto, ON | (613) 879-8085 | ndao298.bkb@gmail.com
resume.daonam.workers.dev | linkedin.com/in/ndao1104 | github.com/ndao1104

EDUCATION

University of Ottawa

Honours BSc in Computer Science

Ottawa, ON

Sep 2017 – May 2022

TECHNICAL SKILLS

Networking & Systems: Linux, Windows Server, TCP/IP, DNS, DHCP, VPN, Firewalls, TLS/SSL, Load Balancing

Monitoring & Logging: Splunk (SPL), Trellix Helix (TQL/MQL), Grafana, NXLog

Cloud & Infrastructure: AWS, Azure, Kubernetes, Docker, Terraform, Ansible, GitLab CI/CD, Mainframe

Automation & Scripting: Python, Bash, PowerShell, REST APIs, Cron

Databases & Tools: SQL, Sybase, DB2, AutoSys, ServiceNow, Git, Jira, Confluence

EXPERIENCE

Security Operations Engineer

Canada Life

Toronto, ON

Oct 2025 – Present

- Provided technical support across multicloud and on-premises environments, troubleshooting application and infrastructure issues with a focus on enterprise security operations.
- Monitored NXLog agent health across 2,000+ servers through SIEM dashboards, investigating telemetry failures and restoring log collection to maintain monitoring coverage.
- Prepared and executed production change requests, coordinating with application, infrastructure, and change-management teams to validate multi-tier environments and support smooth deployments.
- Assessed telemetry across 30+ applications and infrastructure environments, identifying monitoring and observability gaps and addressing approximately 30% of identified coverage gaps before production SIEM onboarding.
- Validated post-deployment log ingestion and parsed fields, developing 50+ custom SPL/MQL queries and documenting troubleshooting and triage procedures in Confluence for use by 25+ analysts.
- Standardized multi-cloud log schemas, evaluating baseline data quality, volume, and ingestion requirements before onboarding to prevent SIEM capacity issues.

Site Reliability Engineer

Morgan Stanley

Montreal, QC

Mar 2023 – Aug 2024

- Provided 24/7 production support for mainframe and distributed trading applications, resolving incidents, performance bottlenecks, and data discrepancies across millions of daily financial transactions.
- Diagnosed Linux production incidents by monitoring Grafana trade queues, analyzing Splunk logs, and debugging Java (Spring) microservices to coordinate rapid cross-functional resolution with development, infrastructure, and business teams.
- Onboarded logs into Splunk and built real-time dashboards to track system performance and trade execution failures. Implemented 50+ automated alerts for proactive monitoring, reducing troubleshooting time from 30 minutes to 5 minutes.
- Built automated data-remediation utilities using Python and SQL to detect and correct data anomalies in Sybase/DB2, reducing manual intervention by 75%.
- Developed and executed ad-hoc SQL queries to resolve trade data discrepancies, enhancing data integrity and compliance with regulatory standards across production systems.
- Managed AutoSys batch jobs in production pipelines, archiving 100M+ trade records weekly and ensuring availability for downstream analytics and compliance.
- Automated infrastructure configuration with Ansible and implemented CI/CD pipelines in GitLab for microservice deployments, reducing environment setup time and improving release frequency.

PROJECTS

Home Server | Linux (Ubuntu), Docker, Docker Compose, Bash, Caddy, Tailscale, WireGuard, Cron, TLS/HTTPS, DNS

- Repurposed an old laptop into an always-on headless Ubuntu home server and configured static IP addressing and DHCP reservations for reliable network connectivity.
- Containerized multiple self-hosted applications (Vaultwarden password manager, AdGuard Home DNS filter) via Docker and Docker Compose, with persistent data volumes and inter-container networking.
- Implemented Caddy as a reverse proxy with automatic HTTPS/TLS and Tailscale for secure remote access to self-hosted services without exposing ports to the public internet, including custom TLS certificates for internal domains.
- Automated encrypted backups with Bash and cron, applying scheduled execution and retention policies.

Ansible Playbook for Self-Hosted Server | Ansible | [Link](#)

- Converted a manually-configured Linux server into Infrastructure as Code using Ansible, writing a playbook to automate Docker installation, system configuration (power management, DNS), container deployment, and backup scheduling.
- Built reusable Jinja2 templates for configuration files (Docker Compose, reverse proxy config, backup scripts), parameterized via Ansible variables to support deployment on different hosts without code changes.
- Used Ansible modules (lineinfile, cron, template, systemd, docker_compose_v2) to manage system state declaratively instead of manual one-off shell commands.

Path of Exile Market Analyzer | Python (Flask), JavaScript, REST APIs | [Link](#)

- Built a full-stack web application in Flask that fetches live game economy data from a third-party public API (poe.ninja) and calculates real-time currency conversions and item flip profitability.
- Developed a backend proxy layer to handle server-side requests to the third-party API, keeping API communication and data processing off the client.
- Designed and implemented interactive UI components in vanilla JavaScript, including a live-updating data table with editable inputs that recalculate results instantly without a page reload, and deployed the app to production.

Flask doc2pdf | Python (Flask), JavaScript, Docker, LibreOffice, Render, CI/CD | [Link](#)

- Built a Flask API that handles file uploads and converts DOCX to PDF using LibreOffice in headless mode via subprocess calls.
- Developed a vanilla JavaScript client with drag-and-drop file uploads using the Fetch API and FormData, supporting sequential multi-file processing with per-file status feedback.
- Containerized the application with Docker, including system-level dependency installation for LibreOffice, and deployed to Render with automated CI/CD builds triggered by GitHub pushes.

Personal Website | HTML/CSS/JS, Git, Cloudflare Pages, CI/CD | [Link](#)

- Designed and built a responsive personal portfolio site from scratch.
- Migrated hosting to Cloudflare Pages and configured CI/CD for automatic builds and deployments on every GitHub push. Set up custom domain routing and HTTPS via Cloudflare's edge network.

PUBLICATION

DistSMOBN: Distributed SMOGN for Imbalanced Regression Problems | [Link](#) – [GitHub](#) – [PDF](#)

- Imbalanced data poses challenges for learning systems, including regression tasks. Existing resampling methods for imbalanced regression often become impractical with large datasets.
- We propose DistSMOBN, the first distributed resampling solution designed for large-scale imbalanced regression.
- DistSMOBN uses Scalable K-Means++ to preserve the spatial relationships of rare cases, followed by SMOGN within each cluster to generate high-quality synthetic samples. Implemented using MapReduce and Apache Spark, our experiments demonstrate its effectiveness and scalability.
- The source code is freely available on GitHub.